

# Neural Networks practical

*Marc Sabaté*

*08/03/2019*

## Deep Learning practical in R

In this practical we will build and train a feedforward network with sigmoid activated neurons from scratch.

### Load the necessary libraries

We load ggplot2 for plotting purposes. If ggplot2 is not installed, run first `install.packages('ggplot2')`.

```
library(ggplot2)
```

### Data loading and visualisation

Create a directory *DL\_practical* in your home directory and save the file *data.txt*. We read the data:

```
HOME = Sys.getenv(x='HOME')
setwd(paste(HOME,"DL_practical",sep="/"))
data = read.table("data.txt", sep=",", header = TRUE)
```

We explore the data. As you can see, each observation of the dataset is two-dimensional with a binary output. The goal of this practical is to fit a binary classifier on the data.

```
summary(data)
```


```
##           x1           x2           y
##  Min.    :-1.4637  Min.    :-0.9431  Min.    :0.0
## 1st Qu.: -0.1236  1st Qu.: -0.1601  1st Qu.: 0.0
## Median :  0.4607  Median :  0.2435  Median : 0.5
## Mean    :  0.5003  Mean    :  0.2561  Mean    : 0.5
## 3rd Qu.:  1.1205  3rd Qu.:  0.6904  3rd Qu.: 1.0
## Max.    :  2.5208  Max.    :  1.4685  Max.    : 1.0
```

```
table(data$y)
```

```
##
##  0  1
## 250 250
```

Execute the following lines to visualise it.

```
p = ggplot(data=data, aes(x1,x2, color=factor(y))) + geom_point()
print(p)
```



## Logistic Regression

As a first choice, we try to fit a Logistic regression on the data. For this, we use the R function `glm()` taking into account that its parameter `family` providing the error distribution should be `'binomial'`.

```
model = glm(formula = y~x1+x2, data = data, family = "binomial")
```


We create a grid in order to test our model. The output of the Logistic regression is a numeric value between 0 and 1, providing the probability of the observation being 1. In other words,  $f_{\theta}(x_1, x_2) = P(Y = 1|x_1, x_2; \theta)$ . Therefore, we create the binary variable `y_grid_binary`, that outputs the class 1 when the probability is greater than 0.5.

```
grid = expand.grid(x1 = seq(min(data$x1), max(data$x1), 0.02),
                  x2 = seq(min(data$x2), max(data$x2), 0.02))

y_grid = predict(model, grid, type="response")
y_grid_binary = as.numeric(y_grid>0.5)
grid$y = y_grid_binary
```

We plot the prediction of the model, along with the training set, in order to see how well the logistic regression is doing.

```
p = ggplot(data=data, aes(x1,x2, color=factor(y))) + geom_point(size=1)
p = p + geom_point(data=grid, aes(x1, x2, color=factor(y)), alpha=0.1)
print(p)
```



As you can see, the Logistic Regression is trying to linearly separate the training set, and therefore it doesn't work well in those datasets that are not linearly separable.

## Neural Networks

### Ingredients

- Definition of the model to do a forward pass:  $y = \sigma(W^{[L-1]}(\sigma(W^{[L-2]}(\dots\sigma(W^{[1]}x + b^{[1]})\dots) + b^{[L-2]})) + b^{[L-1]}$  where
  - $\sigma$  is the activation function
  - $Wx + b$  is a linear function that goes from  $\mathbb{R}^n$  to  $\mathbb{R}^m$ .
  - In our example, the output of the model is a value between 0 and 1 that tells the probability of a point being blue or red. The input of the model are the coordinates of the point.
- Definition of a cost function that tells how good is the model in terms of its parameters:

$$J(\theta), \text{ where } \theta := \{W^{[1]}, b^{[1]}, \dots, W^{[L-1]}, b^{[L-1]}\}$$

- Optimisation of the cost function using gradient descent.
  - We need to apply the chain rule (backpropagation) in order to obtain  $\partial_{\theta} J$  for each optimisation step.

### Activation function and model definition

We will use the `sigmoid` activation function. Please create a function that returns the output of the sigmoid,  $\sigma(z)$  where  $\sigma(z) = \frac{1}{1+\exp(-z)}$

```
# activation function
sigmoid <- function(z){
  # TODO
  #output =
  #return(output)
}
```

We will create a model with a two-dimensional input layer, one hidden layer, and a one-dimensional output layer.

$$y = \sigma(W^{[2]}(\sigma(W^{[1]}x + b^{[1]})) + b^{[2]})$$

We define the next function that goes forward through the model. It takes the input  $\mathbf{x}$ , and return the output of the model, as well as the intermediate layers that we will use in backpropagation.

Please fill in the gaps below taking into account that:  $z_2 = W^{[1]}x + b^{[1]}$ ,  $a_2 = \sigma(z_2)$ ,  $z_3 = W^{[2]}a_2 + b^{[2]}$ ,  $a_3 = \sigma(z_3)$ . For this, take into account that in R matrix multiplication is performed using `%*%`, and the transposed of a matrix  $\mathbf{x}$  is performed using `t(x)`. You will only need to tranpose the input  $\mathbf{x}$  of the model, as in our data the observations are given per row, and we want to put them into columns.

```
# forward model
forward_model <- function(model, x){
  # model is a list with the following
  # x is the input
  W1 = model[['W1']]
  b1 = model[['b1']]
  W2 = model[['W2']]
  b2 = model[['b2']]

  # forward pass
  # hidden layer
  #z2 = #TODO
  #a2 = #TODO

  # output layer
  #z3 = #TODO
  #a3 = #TODO
  output = list(z2 = t(z2), a2=t(a2), z3=t(z2), a3=t(a3))

  return(output)
}
```

## Loss function

Let  $\mathcal{D} = \{(x^1, y^1), (x^2, y^2), \dots, (x^N, y^N)\}$  be our training set, where  $x^i \in \mathbb{R}^n$ . We define the loss function as

$$J(\theta) = - \sum_{i=1}^N y^i \log(f_{\theta}(x^i)) + (1 - y^i) \log(1 - f_{\theta}(x^i))$$

, which corresponds to the negative log-likelihood of of our dataset assuming it was sampled from a Bernouilli distribution.

```
loss_fn <- function(model, x, y){
  layers_model = forward_model(model, x)
  y_pred = layers_model[['a3']]
  loss = y * log(y_pred) + (1-y) * log(1-y_pred)
```

```

    return(sum(-loss))
}

```

### Gradient descent step.

See for example (<https://arxiv.org/pdf/1801.05894.pdf>) for a general form of the gradient of the cost function for the parameters in each layer of the network. The following function performs a gradient descent step using the whole training set

$$\theta := \theta - \alpha \cdot \partial_{\theta} J(\theta)$$

```

GD_step <- function(model, x, y, lr=0.001){
  W1 = model[['W1']]
  b1 = model[['b1']]
  W2 = model[['W2']]
  b2 = model[['b2']]

  pred_model = forward_model(model, x)
  z2 = pred_model[['z2']]
  a2 = pred_model[['a2']]
  z3 = pred_model[['z3']]
  a3 = pred_model[['a3']]

  delta3 = a3 - y
  dW2 = t(a2) %*% delta3
  db2 = apply(FUN=sum, X=delta3, MARGIN=2)

  delta2 = sigmoid(z2) * (1-sigmoid(z2)) * (delta3%*%W2)
  dW1 = t(x) %*% delta2
  db1 = apply(FUN=sum, X=delta2, MARGIN=2)

  W2 = W2 - lr * t(dW2)
  b2 = b2 - lr * db2
  W1 = W1 - lr * t(dW1)
  b1 = b1 - lr * db1

  model[['W1']] = W1
  model[['b1']] = b1
  model[['W2']] = W2
  model[['b2']] = b2

  return(model)
}

```

We put everything together for the training:

- Initialise  $W^{[1]}, b^{[1]}, W^{[2]}, b^{[2]}$ .
- While Not convergence:
  - Calculate  $J(\theta)$
  - Update  $W^{[i]} := W^{[i]} - \alpha \cdot \partial_{W^{[i]}} J$
  - Update  $b^{[i]} := b^{[i]} - \alpha \cdot \partial_{b^{[i]}} J$

```

train <- function(model, n_epochs, x, y){
  for(epoch in 1:n_epochs){

```

```

    model = GD_step(model, x=x, y=y)
    loss = loss_fn(model, x=x, y=y)

    if(epoch %% 10 == 0){
      log = paste("Epoch: ", epoch, "/", n_epochs, ", loss: ", loss, sep="")
      print(log)
    }
  }
  return(model)
}

n_hidden = 30
matrix
W1 = matrix(data=rnorm(2*n_hidden), nrow=n_hidden, ncol=2)
b1 = rnorm(n_hidden)
W2 = matrix(data=rnorm(n_hidden), nrow=1, ncol=n_hidden)
b2 = rnorm(1)
model = list(W1 = W1, b1=b1, W2=W2, b2=b2)

model = train(model=model, n_epochs=10000, x=as.matrix(x=data[,c("x1", "x2")] ), y=c(x=data$y))

We plot the results with the new model.

grid = expand.grid(x1 = seq(min(data$x1), max(data$x1), 0.02),
                  x2 = seq(min(data$x2), max(data$x2), 0.02))
y_grid = forward_model(model, x=as.matrix(grid))[['a3']]
y_grid_binary = y_grid>0.5
grid$y = as.numeric(y_grid_binary)

p = ggplot(data=data, aes(x1,x2, color=factor(y))) + geom_point(size=1)
p = p + geom_point(data=grid, aes(x1, x2, color=factor(y)), alpha=0.1)
print(p)

```